

Data Access Object

Also Known As

DAO

Brief Description

Code that depends on specific features of data resources ties together business logic with data access logic. This makes it difficult to replace or modify an application's data resources.

The **Data Access Object** (or **DAO**) pattern:

- separates a data resource's client interface from its data access mechanisms
- adapts a specific data resource's access API to a generic client interface

The DAO pattern allows data access mechanisms to change independently of the code that uses the data.

Detailed Description

See the [Core J2EE™ Patterns](#)

Detailed Example

The Java Pet Store sample application uses the DAO pattern both for database vendor-neutral data access, and to represent XML data sources as objects.

- **Accessing a database with a DAO.**

A Data Access Object class can provide access to a particular data resource without coupling the resource's API to the business logic. For example, sample application classes access catalog categories, products, and items using DAO interface `CatalogDAO`

Reimplementing `CatalogDAO` for a different data access mechanism (to use a Connector, for example), would have little or no impact on any classes that use `CatalogDAO` because only the implementation would change. Each potential alternate implementation of `CatalogDAO` would access data for the items in the catalog in its own way, while presenting the same API to the class that uses it.

The following code excerpts illustrate how the sample application uses the DAO pattern to separate business logic from data resource access mechanisms:

- Interface [CatalogDAO](#) defines the DAO API. Notice that the methods in the interface below make no reference to a specific data access mechanism. For example, none of the methods specify an SQL query, and they throw only exceptions of type [CatalogDAOException](#) . Avoiding mechanism-specific information in the DAO interface, including exceptions thrown, is essential for hiding implementation details.

```
public interface CatalogDAO {
    public Category getCategory(String categoryID, Locale l)
        throws CatalogDAOException;
    public Page getCategories(int start, int count, Locale l)
        throws CatalogDAOException;
    public Product getProduct(String productID, Locale l)
        throws CatalogDAOException;
    public Page getProducts(String categoryID, int start, int count, Locale l)
        throws CatalogDAOException;
```

```

    public Item getItem(String itemID, Locale l)
        throws CatalogDAException;
    public Page.getItems(String productID, int start, int size, Locale l)
        throws CatalogDAException;
    public Page.searchItems(String query, int start, int size, Locale l)
        throws CatalogDAException;
}

```

- Class [CloudscapeCatalogDAO](#) implements this interface for the Cloudscape relational database, as shown in the following code excerpt. Note that the SQL to access the Cloudscape database is hard-coded.

```

public class CloudscapeCatalogDAO implements CatalogDAO {
    ...
    public static String GET_CATEGORY_STATEMENT
        = "select name, descn "
        + " from (category a join category_details b on a.catid=b.catid) "
        + " where locale = ? and a.catid = ?";
    ...
    public Category getCategory(String categoryID, Locale l)
        throws CatalogDAException {
        Connection c = null;
        PreparedStatement ps = null;
        ResultSet rs = null;
        Category ret = null;

        try {
            c = getDataSource().getConnection();
            ps = c.prepareStatement(GET_CATEGORY_STATEMENT,
                                    ResultSet.TYPE_SCROLL_INSENSITIVE,
                                    ResultSet.CONCUR_READ_ONLY);

            ps.setString(1, l.toString());
            ps.setString(2, categoryID);
            rs = ps.executeQuery();
            if (rs.first()) {
                ret = new Category(categoryID, rs.getString(1), rs.getString(2));
            }

            rs.close();
            ps.close();
            c.close();
            return ret;
        } catch (SQLException se) {
            throw new CatalogDAException("SQLException: " + se.getMessage());
        }
    }
    ...
}

```

- **Implementation strategies.** Designing a DAO interface and implementation is a tradeoff between simplicity and flexibility. The sample application provides examples of several strategies for implementing the Data Access Object pattern.

- *Implement the interface directly as a class.* The simplest (but least flexible) way to implement a data access object is to write it as a class. Class `screenDefinitionsDAO` described below and shown in Figure 1 above, is an example of a class that directly implements a DAO interface. This approach separates

the data access interface from the details of how it is implemented, providing the benefits of the DAO pattern. The data access mechanism can be changed easily by writing a new class that implements the same interface, and changing client code to use the new class. Yet this approach is inflexible because it requires a code changes to modify the data access mechanism.

- *Improve flexibility by making DAOs "pluggable".* A pluggable DAO allows an application developer or deployer to select a data access mechanism with no changes to program code. In this approach, the developer accesses a data source only in terms of an abstract DAO interface. Each DAO interface has one or more concrete classes that implement that interface for a particular type of data source. The application uses a factory object to select the DAO implementation at runtime, based on configuration information.

For example, the sample application uses factory class [CatalogDAOFactory](#) to select the class that implements the DAO interface for the catalog. Figure 2 below presents a structure diagram of the Data Access Object design pattern using a factory to select a DAO implementation.

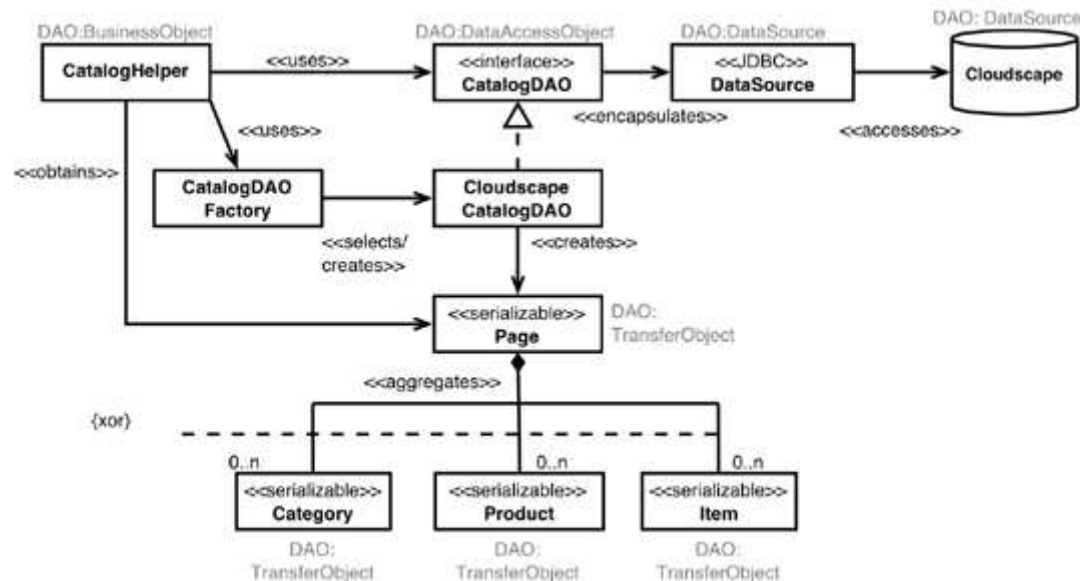


Figure 2. A pluggable DAO

At runtime, the [CatalogHelper](#) uses the [CatalogDAOFactory](#) to create an object that implements [CatalogDAO](#). The factory looks up the name of the class that implements the DAO interface in environment entry "param/CatalogDAOClass". The [CatalogHelper](#) accesses the catalog data source exclusively using the object created by the factory. In the example shown in the figure, the environment entry was set to the (fully-specified) name of class [CloudscapeCatalogDAO](#). This class implements the catalog DAO interface in terms of JDBC™ data sources, accessing a Cloudscape relational database.

This approach is more flexible than using a hard-coded class. To add a new type of data source, an application developer would simply create a class that implements [CatalogDAO](#) in terms of the new data source type, specify the implementing class's name in the environment entry, and re-deploy. The factory would create an instance of the new DAO class, and the application would use the new data source type.

- *Reduce redundancy by externalizing SQL.* Writing a separate class for data source types that have similar APIs can create a great deal of redundant code. For example, JDBC data sources differ from one another primarily in the SQL used to access them. The only differences between the Cloudscape DAO described above and the DAO for a different SQL database are the connection string and the SQL used to access the database.

The sample application reduces redundant code by using a "generic DAO" that externalizes the SQL for different JDBC data sources. Figure 3 below shows how the sample application uses an XML file to specify the SQL for different JDBC data sources.

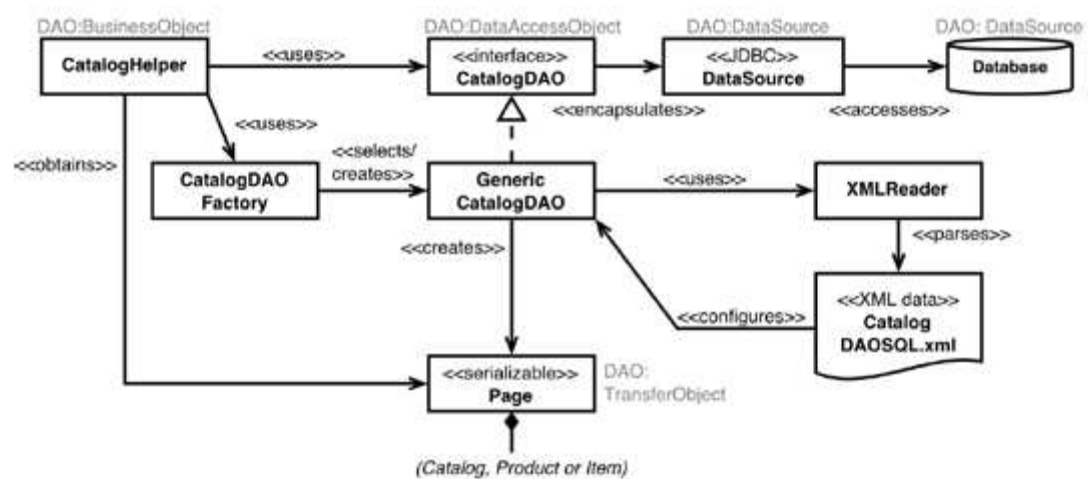


Figure 3. Externalizing DAO SQL

In the figure, the `CatalogDAOFactory` has selected an object of type `GenericCatalogDAO` as the DAO to access the catalog. An XML file called `catalog DAO SQL.xml` specifies the SQL for each supported operation on each type of database. `GenericCatalogDAO` configures itself by selecting a group of SQL statements from the XML file that correspond to the database type named by environment entry "`param/catalog DAO database`". The code sample below shows the definition of the XML for the `getCategory` operation of the `CatalogDAO`. Different SQL is specified for "cloudscape" and "oracle" database types.

```

<DAOConfiguration>
  <SQLStatements database="cloudscape">
    <SQLStatement method="GET_CATEGORY">
      <SQLFragment parameterNb="2">
        select name, descn
          from (category a join category_details b
                on a.catid=b.catid)
         where locale = ? and a.catid = ?
      </SQLFragment>
    </SQLStatement>
    ...
  </SQLStatements>
  <SQLStatements database="oracle">
    <SQLStatement method="GET_CATEGORY">
      <SQLFragment parameterNb="2">
        select name, descn
          from category a, category_details b
         where a.catid = b.catid and locale = ? and a.catid = ?
      </SQLFragment>
    </SQLStatement>
    ...
  </SQLStatements>
</DAOConfiguration>

```

Method `GenericCatalogDAO getCategory` chooses the SQL corresponding to the configured database type, and uses it to fetch a category from the database via JDBC. The following code excerpt shows the implementation of the method.

```

public Category getCategory(String categoryId, Locale locale)
    throws CatalogDAOException {
    Connection connection = null;
    ResultSet resultSet = null;
    PreparedStatement statement = null;
    try {
        connection = getDataSource().getConnection();
        String [] parameterValues = new String [] { locale.toString(), categoryId };
        statement = buildSQLStatement(connection, sqlStatements,
            XML_GET_CATEGORY, parameterValues);
        resultSet = statement.executeQuery();
    }
}

```

```

        if (resultSet.first()) {
            return new Category(categoryID, resultSet.getString(1), resultSet.getString(2));
        }
        return null;
    } catch (SQLException exception) {
        throw new CatalogDAOException("SQLException: " + exception.getMessage());
    } finally {
        closeAll(connection, statement, resultSet);
    }
}
}

```

Notice that the method catches any possible `SQLException` and converts it to a `CatalogDAOException`, hiding the implementation detail that the DAO uses a JDBC database.

This strategy supports multiple JDBC databases with a single DAO class. It both decreases redundant code, and makes new database types easier to add. To support a new database type, a developer simply adds the SQL statements for that database type to the XML file, updates the environment entry to use the new type, and redeploys.

The pluggable DAO and generic DAO strategies can be used separately. If you know that a DAO class will only ever use JDBC databases (for example), the generic DAO class can be hardwired into the application, instead of selected by a factory. For maximum flexibility, the sample application uses both a factory method and a generic DAO.

- **Encapsulating non-database data resources as DAO classes.**

A data access object can represent data that is not stored in a database. The sample application uses the DAO pattern to represent XML data sources as objects. Sample application screens are defined in an XML file which is interpreted by the class `ScreenDefinitionsDAO`. Specifying screen definitions externally makes access to the screen definitions more flexible. For example, if the application designers (or maintainers) decide to change the application to store screen descriptions in the database, instead of in an XML file, they would need only to implement a single new class (`ScreenDefinitionsDAODatabase` for example). The code that uses `ScreenDefinitionsDAO` would remain unchanged, but the data would come from the database via the new class.

The screen definitions mechanism in the sample application provides an example of a concrete Data Access Object representing an underlying, non-database resource (an XML file).

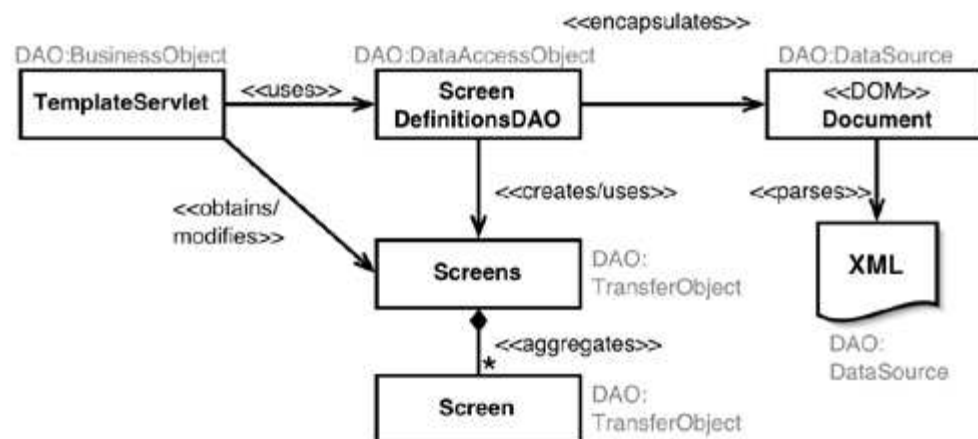


Figure 1. Data Access Object providing access to XML data source

Figure 1 shows a structure diagram of the `ScreenDefinitionsDAO` managing the loading and interpretation of XML data that defines application screens.

- `TemplateServlet` uses the `ScreenDefinitionsDAO` to load screen definitions:

```
Screens screenDefinitions =
```

- [ScreenDefinitionDAO](#) represents screen definitions in an XML file deployed with an application. It uses XML APIs to read the screen definitions from the XML file. Only this class would need to be replaced to support storing screen definitions in some other way. The method that loads screen definitions looks like this:

```

public static Screens loadScreenDefinitions(URL location) {
    Element root = loadDocument(location);
    if (root != null) return getScreens(root);
    else return null;
}
...
public static Screens getScreens(Element root) {
    // get the template
    String defaultTemplate = getTagValue(root, DEFAULT_TEMPLATE);
    if (defaultTemplate == null) {
        System.err.println("*** ScreenDefinitionDAO error: " +
            " Default Template not Defined.");
        return null;
    }

    Screens screens = new Screens(defaultTemplate);
    getTemplates(root, screens);
    // get screens
    NodeList list = root.getElementsByTagName(SCREEN);
    for (int loop = 0; loop < list.getLength(); loop++) {
        Node node = list.item(loop);
        if ((node != null) && node instanceof Element) {
            String templateName = ((Element)node).getAttribute(TEMPLATE);
            String screenName = ((Element)node).getAttribute(NAME);
            HashMap parameters = getParameters(node);
            Screen screen = new Screen(screenName, templateName, parameters);
            if (!screens.containsScreen(screenName)) {
                screens.addScreen(screenName, screen);
            } else {
                System.err.println("*** Non Fatal error: Screen " + screenName +
                    " defined more than once in screen definitions file");
            }
        }
    }
    return screens;
}
...

```

The code fragment above shows how `loadScreenDefinitions` loads screen definitions using DOM interfaces, while hiding that fact from clients of the class. A client of this class can expect to receive a `screens` object regardless of how those screens are loaded from persistent storage. Method `getScreens` handles all of the DOM-specific details of loading a screen from an XML file.